

Progressively include weaker heatmap responses; stop at first predicate flip

Input image



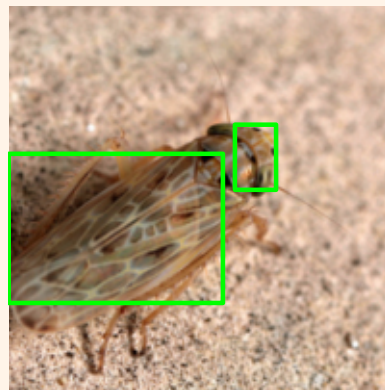
$$p_{(1349, -)}(x) = 1$$

Neuron-targeted heatmap



gradient-based; perturbation-based
fallback

Candidate box union B



connected regions

Targeted blur removal



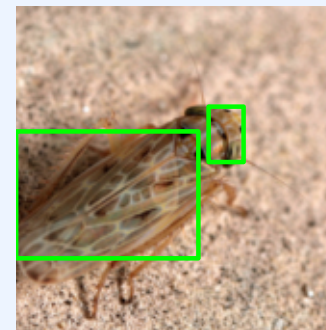
$$p_{(1349, -)}(x_{\setminus B}) = 0$$

predicate deactivated → accept B

Attribution-guided box validation

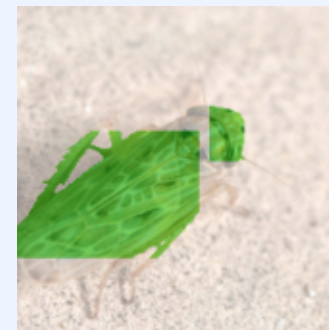
Foreground segmentation mask F

Accepted region



validated: B

Foreground overlay



display only: $B \cap F$

The removal test validates the box.
Segmentation clarifies its visible content.

Foreground-assisted visualization